

---

# Generative models based on ODE and SDE

---

Denis Rakitin  
rakitinden32@gmail.com

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Diffusion models</b>	<b>2</b>
2.1	Score-based models . . . . .	2
2.1.1	Score function . . . . .	2
2.1.2	Denoising score matching . . . . .	3
2.1.3	Noise Conditional Score Networks . . . . .	5
2.2	Preliminaries on ODEs and SDEs . . . . .	6
2.2.1	Ordinary differential equations . . . . .	6
2.2.2	Stochastic differential equations: informal . . . . .	7
2.2.3	Continuity equation . . . . .	9
2.2.4	Fokker-Planck equation . . . . .	11
2.2.5	Uniqueness and physical interpretation . . . . .	12
2.3	Diffusion models based on SDEs . . . . .	14
2.3.1	Practical usage . . . . .	15
2.3.2	Calculating likelihood . . . . .	16
2.4	Conditional diffusion models . . . . .	18
2.4.1	Conditional training . . . . .	18
2.4.2	Classifier guidance . . . . .	18
2.4.3	Classifier-free guidance . . . . .	19

# 1 Introduction




Figure 1: Top: denoising process, defined by a diffusion model [25]. Bottom: stochastic interpolation between two pictures [1].

This tutorial aims to introduce the reader to the mathematical methods which are widely used in contemporary generative modeling. Big practical success of diffusion models [14, 8, 25], which generate a sequence of pictures instead of just the target one, led to development of a novel family of models that describe some dynamic processes. Almost all processes occurring in the real world can be described by differential equations, which will be the basis of the generative models, reviewed here. We will cover methods based on either ordinary (ODE) or stochastic (SDE) differential equations, applicable for generation [25, 19, 26, 2, 1], paired [26, 2, 1, 20] and unpaired [21, 23, 18, 13] domain translation, formalized as an instance of the optimal transport problem.

## 2 Diffusion models

Generally speaking, diffusion models define a process of step-by-step noising of a picture and try to learn the reverse process, which allows to generate new images starting with a pure noise. There are different approaches for formalizing this concept. The earlier score-based models as NCSN [24] learn score functions of probability distributions at all noise levels and sample with a consequent Langevin dynamics [27] going from larger to smaller noise levels. In [14] the backward denoising process is trained as a latent variable model with a variational distribution on noisy images, corresponding to a forward noising process. Finally, the authors of [25] consider a continuous-time noising SDE and construct the reverse, which can be seen as a unifying framework for previous approaches. On the one hand, the backward SDE explicitly requires knowing score functions at all noise levels, which is similar to the earlier score-based models. On the other hand, discrete-time sampling schemes from the backward SDE result in a discrete-time process very similar to the one in [14]. In this section, we will cover the original score-based approach and models, based on SDE.

### 2.1 Score-based models

The general problem of generative modeling consists of constructing the probabilistic model  $p_\theta$  given a data set  $X_1, \dots, X_n \sim p_{data}$  in such a way that  $p_\theta \approx p_{data}$ .

There are mainly 3 families of non-diffusion generative models: Generative Adversarial Networks (GANs) [11], Variational Autoencoders (VAEs) [17] and Normalizing Flows (NFs) [10, 22]. Despite different training procedures, all of them share the same simple generation scheme: picture  $x$  is obtained by transforming noise  $z$  by a trained network  $G_\theta$ :  $x = G_\theta(z)$ . In contrast, score-based models do not learn the generator explicitly, but try to approximate the score function of the distribution:  $s_\theta(x) \approx \nabla_x \log p_{data}(x)$ . But why is this statistic valuable?

#### 2.1.1 Score function

First, it can be used for finding mode of the distribution by gradient ascent: if

$$X_{t+1} = X_t + \gamma \nabla \log p_{data}(X_t),$$

then

$$X_t \xrightarrow[t \rightarrow \infty]{} x^* = \arg \max_x \log p_{data}(x) = \arg \max_x p_{data}(x).$$

Of course, in general multi-modal case method converges to one of the stationary points of  $\log p_{data}(x)$ .

Furthermore, it can be used to obtain samples from the distribution using Langevin dynamics [27]:

$$X_{t+1} = X_t + \gamma \nabla \log p_{data}(X_t) + \sqrt{2\gamma} \varepsilon_t, \quad (1)$$

where  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent with  $X_t$ . In regular cases, one can prove that

$$p_{X_t} \xrightarrow[t \rightarrow \infty]{} p^\gamma,$$

where  $p^\gamma$  is a distribution close to  $p_{data}$  that depends on a discretization step and converges to  $p_{data}$ , when  $\gamma$  converges to zero:

$$p^\gamma \xrightarrow[\gamma \rightarrow 0]{} p_{data}.$$

Given this, score-based models obtain sample from the model by running Langevin dynamics with trained approximation of the score function  $s_\theta$  instead of the true one:

$$X_{t+1} = X_t + \gamma s_\theta(X_t) + \sqrt{2\gamma} \varepsilon_t.$$

The only question remained is how to train it. Ideally, we would like to match our approximation with the true score function on the samples from data set and solve the Score Matching objective

$$\mathbb{E} \|s_\theta(X) - \nabla \log p_{data}(X)\|^2 \rightarrow \min_\theta, \quad (2)$$

where  $X \sim p_{data}$ . Unfortunately, this objective has 2 issues:

1. Trivially, we do not know the true score function and cannot compute it to perform regression;
2. It is common to assume that such high-dimensional structured data like images lie in a much less-dimensional manifold. In this case, data distribution  $p_{data}$  does not have density in a regular sense. Even if the data distribution is not strictly concentrated on a manifold, but is very close, there will be very rapid transitions from zero to high-density regions, which will result in large unstable values of  $\nabla \log p_{data}$ .

### 2.1.2 Denoising score matching

To address the second problem, one can make the data distribution more smooth by, for example, adding gaussian noise:

$$\hat{X} = X + \sigma \varepsilon, \varepsilon \sim \mathcal{N}(0, I),$$

or, more generally, deleting part of the object and replacing it with gaussian noise:

$$\hat{X} = \alpha X + \sigma \varepsilon, \varepsilon \sim \mathcal{N}(0, I), \alpha < 1.$$

In terms of distributions, we defined a conditional distribution

$$p_{\hat{X}|X}(\hat{x}|x) = \mathcal{N}(\hat{x} | \alpha x, \sigma^2 I).$$

The perturbed variable  $\hat{X}$  always has density, which can be represented as

$$p_{\hat{X}}(\hat{x}) = \int p_{\hat{X}|X}(\hat{x}|x) p_X(x) dx. \quad (3)$$

Consequently, it has a score function  $\nabla \log p_{\hat{X}}(\hat{x})$ , which can be theoretically learned by optimizing the score matching objective

$$\mathbb{E} \|s_\theta(\hat{X}) - \nabla \log p_{\hat{X}}(\hat{X})\|^2 \rightarrow \min_\theta,$$

which is, however, still intractable.

The representation of density in the Equation 3 is important, because it rewrites a complicated density  $p_{\hat{X}}$  as an integral of a very easy conditional density  $p_{\hat{X}|X}$  with respect to the distribution  $p_X$ ,

from which we have a data set of samples. It turns out that the score function of  $p_{\hat{X}}$  has a similar representation. To derive it, we start with differentiating logarithm (so-called log-derivative trick):

$$\nabla_{\hat{x}} \log p_{\hat{X}}(\hat{x}) = \frac{\nabla_{\hat{x}} p_{\hat{X}}(\hat{x})}{p_{\hat{X}}(\hat{x})}.$$

Next, we use the representation of density  $p_{\hat{X}}(\hat{x})$  in Eq. 3 and obtain

$$\frac{\nabla_{\hat{x}} \int p_{\hat{X}|X}(\hat{x}|x) p_X(x) dx}{p_{\hat{X}}(\hat{x})} = \frac{\int \nabla_{\hat{x}} p_{\hat{X}|X}(\hat{x}|x) p_X(x) dx}{p_{\hat{X}}(\hat{x})}.$$

Finally, we apply the reversed log-derivative trick to  $\nabla_{\hat{x}} p_{\hat{X}|X}(\hat{x}|x)$  and obtain

$$\frac{\int \nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{x}|x) \cdot p_{\hat{X}|X}(\hat{x}|x) p_X(x) dx}{p_{\hat{X}}(\hat{x})} = \int \nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{x}|x) \cdot p_{X|\hat{X}}(x|\hat{x}) dx,$$

which is just the conditional expectation of the conditional score function. Thereby, we obtained:

$$\nabla_{\hat{x}} \log p_{\hat{X}}(\hat{x}) = \mathbb{E} \left[ \nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{X}|X) \mid \hat{X} = \hat{x} \right]. \quad (4)$$

Given this, the score matching objective can be rewritten as

$$\mathbb{E} \left\| s_{\theta}(\hat{X}) - \mathbb{E} \left[ \nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{X}|X) \mid \hat{X} \right] \right\|^2 \rightarrow \min_{\theta}. \quad (5)$$

Conditional expectation is very convenient to work with since it is the best predictor in the  $L_2$  sense: for all pairs of r.v.  $(X, Y)$

$$g^*(x) = \mathbb{E}[Y \mid X = x] = \arg \min_g \mathbb{E} \|g(X) - \mathbb{E}[Y|X]\|^2 = \arg \min_g \mathbb{E} \|g(X) - Y\|^2,$$

which means that the objective in Eq. 5 is equivalent to the so-called «denoising score matching» objective

$$\mathbb{E} \|s_{\theta}(\hat{X}) - \nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{X}|X)\|^2 \rightarrow \min_{\theta}. \quad (6)$$

Surprisingly, by addressing the problem of sharp density transitions, we also implicitly solved the intractability of the objective! The conditional score function is just the score function of the normal distribution, which can be calculated. For  $d$ -dimensional normal distribution

$$p_{\hat{X}|X}(\hat{x}|x) = \mathcal{N}(\hat{x}|\alpha x, \sigma^2 I) = \frac{1}{(2\pi\sigma^2)^{d/2}} \exp \left( -\frac{1}{2\sigma^2} \|\hat{x} - \alpha x\|^2 \right)$$

log-density is equal to

$$\log p_{\hat{X}|X}(\hat{x}|x) = \text{const} - \frac{1}{2\sigma^2} \|\hat{x} - \alpha x\|^2,$$

which gives the score function

$$\nabla_{\hat{x}} \log p_{\hat{X}|X}(\hat{x}|x) = -\frac{1}{\sigma^2} (\hat{x} - \alpha x) = \frac{1}{\sigma^2} (\alpha x - \hat{x}).$$

Finally, we arrive at optimizing the objective

$$\mathbb{E} \left\| s_{\theta}(\hat{X}) - \frac{1}{\sigma^2} (\alpha X - \hat{X}) \right\|^2 \rightarrow \min_{\theta},$$

which minimal value is obtained at the score function of the perturbed data distribution:

$$s_{\theta^*}(\hat{x}) = \nabla_{\hat{x}} \log p_{\hat{X}}(\hat{x}) = \nabla_{\hat{x}} \log \int p_{data}(x) \mathcal{N}(\hat{x}|\alpha x, \sigma^2 I) dx.$$

To sample from a trained score-based model, one can just apply Langevin dynamics, which is summarized in the Algorithm 1.

Formally, we obtained algorithm which allows to sample from a slightly modified data distribution. In theory, one can obtain a very close approximation of  $p_{data}$  by setting  $\alpha \approx 1$  and  $\sigma^2 \approx 0$ . In practice, however, one should keep in mind that when they get close to these values, transitions of density become sharp, score function starts to take large values and the training procedure becomes unstable. This leads to a trade-off between precision of the approximation and stability of the model.

---

**Algorithm 1** Sampling from a score-based model

---

```

 $X^0 \sim p_0$ 
for  $m = 1, \dots, M$  do
   $\varepsilon^{(m)} \sim \mathcal{N}(0, I)$  — independent with  $X^{(m-1)}$ 
   $X^{(m)} = X^{(m-1)} + \gamma_m s_\theta(X^{(m-1)}) + \sqrt{2\gamma_m} \varepsilon^{(m)}$  ▷ Langevin step
end for
return  $X^{(M)}$ 

```

---

### 2.1.3 Noise Conditional Score Networks

To overcome the necessity of choosing  $\alpha$  and  $\sigma^2$  and balancing between two qualities of the model, authors of [24] present a very elegant idea: consider a sequence of modified variables  $\{X_t\}_{t=1}^T$  instead of one  $\hat{X}$ . This sequence of variables will represent a process of gradual noising of the image and cover the whole spectrum of noisy distributions: from sharp distributions close to  $p_{data}$  to the pure noise like  $\mathcal{N}(0, I)$ . Formally, this sequence should possess 3 qualities:

1. The first variable  $X_1$  should be close to the data distribution:  $p_{X_1} \approx p_{data}$ . This will ensure that sampling from  $p_{X_1}$  will be almost equivalent to sampling from  $p_{data}$ ;
2. The last variable  $X_T$  should be a very simple distribution to sample from, for example,  $\mathcal{N}(0, \sigma^2)$ ;
3. Distributions of  $X_t$  and  $X_{t+1}$  should be close to each other. This will make  $X_{t+1}$  a good initialization point for sampling from  $p_{X_t}$ .

Below, we will use notation of the form  $p_t(x_t) := p_{X_t}(x_t)$  or  $p_{t|s}(x_t|x_s) := p_{X_t|X_s}(x_t|x_s)$  to make it shorter. Authors consider a so-called variance-exploding (VE) process and define

$$X_t = X_0 + \sigma_t^2 \varepsilon, \quad (7)$$

where  $\varepsilon \sim \mathcal{N}(0, I)$  is independent from  $X_0$  and  $\sigma_t$  is an increasing sequence of variances. Equivalently,

$$p_{t|0}(x_t|x_0) = \mathcal{N}(x_t|x_0, \sigma_t^2 I).$$

Taking  $\sigma_0 \approx 0$ ,  $\sigma_t \approx \sigma_{t+1}$  and  $\sigma_T^2$  of such magnitude, that  $p_T \approx \mathcal{N}(0, \sigma_T^2)$ , one ensures to match all the 3 requirements. This sequence of distributions corresponds to adding more and more noise to the original distribution, until it becomes indistinguishable from the pure noise with large magnitude.

The more popular process now is the variance preserving process, which defines a Markov chain

$$X_{t+1} = \sqrt{1 - \beta_t} X_t + \sqrt{\beta_t} \varepsilon_t, \quad (8)$$

where  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent from  $X_t$ . This defines a conditional distribution given previous time step:

$$p_{t+1|t}(x_{t+1}|x_t) = \mathcal{N}(x_{t+1}|\sqrt{1 - \beta_t}x_t, \beta_t I).$$

Given this, one can calculate the conditional distribution given the initial variable and obtain [14]

$$p_{t|0}(x_t|x_0) = \mathcal{N}(x_t|\alpha_t x_0, \sigma_t^2 I),$$

where  $\alpha_t = \sqrt{\prod_{s=1}^t (1 - \beta_s)} \rightarrow 0$  and  $\sigma_t^2 = 1 - \prod_{s=1}^t (1 - \beta_s) \rightarrow 1$  given a proper choice of  $\beta_t$ . Choosing small enough  $\beta_t$  to ensure  $p_{t+1} \approx p_t$  and fulfilling  $\alpha_T \approx 0$  and  $\sigma_T^2 \approx 1$ , one will satisfy all the 3 requirements.

Now, given a sequence of modified distributions with easy conditional distributions, one can train score for all of them with denoising score matching:

$$\mathbb{E} \|s_\theta(X_t, t) - \nabla \log p_{t|0}(X_t|X_0)\|^2 \rightarrow \min_{\theta}.$$

In practice, training  $T$  neural networks is very inefficient, that is why  $s_\theta(x, t)$  is defined as one neural network with two inputs. Besides, scores for different  $t$  are connected and the training signal from one

time step is beneficial for another. Thus, the final training procedure consists of taking the weighted sum of denoising score matching losses from all the time steps:

$$\sum_{t=1}^T \gamma(t) \mathbb{E} \|s_\theta(X_t, t) - \nabla \log p_{t|0}(X_t|X_0)\|^2 \rightarrow \min_{\theta}. \quad (9)$$

This model (paired with the sampling Algorithm 2) is called Noise Conditional Score Network [24]. Theoretically,  $s_\theta(x, t)$  matches  $\nabla \log p_t(x)$  after training. Sampling procedure consists of the same Langevin dynamics, but now performed sequentially for all of the time steps backwards. Here the 3 properties of the sequence  $p_t$  become crucial: generating from  $p_T$  is easy, sample from  $p_t$  is a good point to start dynamics for  $p_{t-1}$  and sample from  $p_1$  is almost a sample from  $p_{data}$ . Formally, the sampling procedure is written in the Algorithm 2.

---

**Algorithm 2** Sampling from a Noise Conditional Score Network (NCSN)

---

```

 $X_T^{(M)} \sim p_T$ 
for  $t = T - 1, \dots, 1$  do
   $X_t^{(0)} = X_{t+1}^{(M)}$ 
  for  $m = 1, \dots, M$  do                                      $\triangleright$  Langevin dynamics for  $p_t$ 
     $\varepsilon_t^{(m)} \sim \mathcal{N}(0, I)$  — independent with  $X_t^{(m-1)}$ 
     $X_t^{(m)} = X_t^{(m-1)} + \gamma_m s_\theta(X_t^{(m-1)}, t) + \sqrt{2\gamma_m} \varepsilon_t^{(m)}$             $\triangleright$  Langevin step
  end for
end for
return  $X_1^{(M)}$ 

```

---

We obtained a model with sequential sampling, which needs to use a neural network multiple times. At the same time, it has a very efficient *simulation-free* training procedure, which does not require sampling from the model (and which distinguishes it from energy-based models and continuous normalizing flows). This makes this procedure inefficient in terms of sampling time, but allows to extract a lot of additional information compared to using NN just once. This combination is believed to be very powerful in practice, that is why constructing the analogues for problems other than generation could be very beneficial.

The only major thing that is improved in the newer algorithms is the sampling scheme. It seems like running Langevin dynamics for each time step can be optimized in a way to perform just one step from  $t$  to  $t - 1$ . One of the possible ways to derive such sampling scheme is through stochastic differential equations.

## 2.2 Preliminaries on ODEs and SDEs

### 2.2.1 Ordinary differential equations

Ordinary differential equation, written as

$$\dot{X} = f(X, t),$$

or, equivalently,

$$dX_t = f(X_t, t)dt, \quad (10)$$

defines a set of functions by defining derivative at each point. The most convenient interpretation for us is that  $t$  defines time,  $X_t$  defines position of a particle at time  $t$  and  $f(X_t, t)$  represents its velocity vector at time  $t$  (physical meaning of the derivative). This interpretation is clearly seen from the Euler scheme, that approximates the solution as:

$$X_{t+h} = X_t + h\dot{X}_t + \bar{o}(h) = X_t + hf(X_t, t) + \bar{o}(h) \approx X_t + hf(X_t, t). \quad (11)$$

Euler scheme tells us that position of the particle after small time  $h$  can be obtained by moving particle from the current position  $X_t$  along its velocity field  $f(X_t, t)$  at a distance, proportional to the change in time  $h$ . This approximate solving scheme will allow us to take a first look at SDEs without heavy theory.

Knowing velocity field is not enough to define a unique trajectory of the particle. We also need to know its initial position. The task of solving system

$$\begin{cases} dX_t = f(X_t, t)dt; \\ X_0 = z \end{cases} \quad (12)$$

of the differential equation and the initial value is called Cauchy problem or initial value problem.

By adding the initial condition  $X_0 = z$ , which means that the particle starts moving at the point  $z$ , we fully defined a physical system. Consequently, we defined a unique trajectory on some time segment. Of course, this is not always the case, but some conditions on the velocity  $f(X_t, t)$  guarantee existence and uniqueness of the solution on some time segment  $[0, T]$ . We will not think about it and will assume existence and uniqueness of the solution.

Euler scheme allows to approximate continuous-time processes with discrete-time ones. But this connection also works in the opposite: given a discrete-time process, one can construct its continuous-time generalization, calculate its derivative and represent it as a solution of the corresponding ODE. Let's take, for example, the variance-preserving process, defined in the Eq. 7, and remove its stochastic part. We will obtain a deterministic process of step-by-step deleting of an object (instead of step-by-step noising):

$$X_{t+1} = \sqrt{1 - \beta_t} X_t.$$

The goal is to construct its continuous-time generalization. Let's take now  $t \in [0, 1]$ , small  $h$  and define the connection between  $X_t$  and  $X_{t+h}$  instead of  $X_{t+1}$ . In the original process, coefficient  $\beta_t$  defines portion of object that should be deleted after a unit of time. A logical continuous-time generalization is

$$X_{t+h} = \sqrt{1 - h\beta_t} X_t,$$

which says that the deleted portion of the object is proportional to the time spent. Approximating the square root by Taylor expansion and rearranging terms, we obtain:

$$\frac{X_{t+h} - X_t}{h} = \frac{1}{h} \left( \sqrt{1 - h\beta_t} X_t - X_t \right) = X_t \frac{1 - \frac{h\beta_t}{2} + \bar{o}(h) - 1}{h} = -\frac{\beta_t}{2} X_t + \bar{o}(1).$$

By taking  $h \rightarrow 0$ , we see that the continuous-time analogue can be defined as the solution of the ODE

$$dX_t = -\frac{\beta_t}{2} X_t dt, \quad (13)$$

which is equal to

$$X_t = X_0 \exp \left( -\frac{1}{2} \int_0^t \beta_s ds \right),$$

which is itself a very natural way to define a deleting process, especially, when taking constant  $\beta_t \equiv \beta$ :

$$X_t = X_0 \exp \left( -\frac{\beta t}{2} \right).$$

### 2.2.2 Stochastic differential equations: informal

Stochastic differential equations are far more difficult to define despite their clear physical interpretation. Talking about ODEs, we treated solution of the ODE as a trajectory of the particle that moves under deterministic force. But what if there is also a stochastic force that affects direction of the velocity? This happens, for example, in quantum mechanics, in which some particles move in a purely stochastic way. A natural modification in this case would be to add a stochastic term to the deterministic velocity field and obtain equation of the form

$$dX_t = f(X_t, t)dt + \text{randomness}.$$

In discrete-time one can define a purely stochastic trajectory as a random walk. Let  $X_0 = 0$  and  $X_1, X_2, \dots, X_n, \dots$  be independent variables that define direction of the walk at the corresponding moment:  $P(X_i = 1) = P(X_i = -1) = 1/2$ . Then, position of the random walk at the moment  $n$  is  $S_n = \sum_{i=1}^n X_i$ . This example shows that the «derivative» of a random walk, which in discrete time

can be naturally defined as  $S_{n+1} - S_n$ , is just a set of independent variables with  $\mathbb{E}(S_{n+1} - S_n) = 0$  and  $\text{Var}(S_{n+1} - S_n) = 1$ .

We would like to define a continuous-time analogue of the random walk and its derivative. Given previous observations, the latter seems to be easier to define: let  $(Z_t, t \in [0, 1])$  be a set of independent  $\mathcal{N}(0, I)$  variables, which is called a white noise process. Then one can define an equation of the form

$$dX_t = (f(X_t, t) + Z_t) dt.$$

Where's the problem? One should integrate the white noise  $Z_t$  over time to obtain the trajectory, but  $Z_t$  is nowhere continuous and is not integrable. This motivates to construct the continuous-time random walk first.

**Definition 1** *The continuous-time (d-dimensional) random walk is called Wiener process (or process of Brownian motion), is denoted as  $W_t$ , and possesses 4 properties:*

1.  $W_0 = 0$ ;
2.  $W_t$  is a continuous process;
3. It has independent increments: for all time points  $t_1 < t_2 < \dots < t_n$  variables  $W_{t_1}, W_{t_2} - W_{t_1}, \dots, W_{t_n} - W_{t_{n-1}}$  are independent;
4. Increments  $W_t - W_s$  are normally distributed  $\mathcal{N}(0, (t - s)I)$  variables.

First two properties are clear. Third is the random walk property that tells that the direction of the particle is independent of its current position. Fourth tells that the magnitude of the increments is stationary over time (it only depends on the time difference), which generalizes discrete processes represented as sums of i.i.d. random variables.

Now, given a continuous-time random walk, represented as a Wiener process, one writes a stochastic differential equation as

$$dX_t = f(X_t, t)dt + G(X_t, t)dW_t, \quad (14)$$

where  $G(x, t)$  is in general case matrix, called the *diffusion matrix*. Most of the time, we will use a scalar *diffusion coefficient*  $g(t)$  instead.

But what is the second term formally? The first idea is to say that  $dW_t = W'_t dt$ . However, Wiener process is almost surely nowhere differentiable, and the previous definition has no sense. One should actually treat this equation in an integral form

$$X_t = X_0 + \int_0^t f(X_s, s)ds + \int_0^t G(X_s, s)dW_s \quad (15)$$

and define the latter term (which is called an *Itô integral*) first. However, we will come to this later and now define SDEs by the corresponding discretization scheme. The Euler Scheme, defined in Equation 11, naturally generalizes to the Euler-Maruyama scheme for solving SDEs, which we will use as a pseudo-definition:

$$X_{t+h} \approx X_t + hf(X_t, t) + G(X_t, t)(W_{t+h} - W_t) \quad (16)$$

Recursively applying this scheme up to the initial value  $X_0$ , one can see that  $X_t$  can be represented as a function of increments of the Wiener process before time  $t$ , which are independent with  $W_{t+h} - W_t$ . Given  $W_{t+h} - W_t \sim \mathcal{N}(0, h)$ , one can rewrite scheme as

$$X_{t+h} \approx X_t + hf(X_t, t) + \sqrt{h}G(X_t, t)\varepsilon_t, \quad (17)$$

where  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent with  $X_t$ .

Given a pseudo-definition, let's move on to the examples. In Equation 13 we defined an ODE, corresponding to a deterministic part of the VP process. Now, armed with a «definition» of SDEs, we can finish and construct continuous-time analogue of the entire VP process. As previously, we replace  $X_{t+1}$  with  $X_{t+h}$  and  $\beta_t$  with  $h\beta_t$  and obtain

$$X_{t+h} = \sqrt{1 - h\beta_t}X_t + \sqrt{h\beta_t}\varepsilon_t,$$



where  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent with  $X_t$ . As previously, we use Taylor expansion of the square root and obtain

$$X_{t+h} - X_t = -h \frac{\beta_t}{2} X_t + \bar{o}(h) + \sqrt{h\beta_t} \varepsilon_t \approx -h \frac{\beta_t}{2} X_t + \sqrt{h\beta_t} \varepsilon_t.$$

Comparing it to the Euler-Maruyama scheme, one can see that this is a discretization of the SDE

$$dX_t = -\frac{\beta_t}{2} X_t dt + \sqrt{\beta_t} dW_t, \quad (18)$$

which has many different names: VP-SDE [25], Langevin equation, Ornstein-Uhlenbeck process. Regardless of the name, it defines a continuous-time process of a gradual noising of an object.

But didn't we use any stochastic schemes that look like discretization of an SDE? In Equation 1 we defined Langevin dynamics:

$$X_{t+1} = X_t + \gamma \nabla \log p(X_t) + \sqrt{2\gamma} \varepsilon_t,$$

where  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent with  $X_t$ . Denoting neighbour values of the Langevin dynamics as  $X_t$  and  $X_{t+h}$  and taking step size  $\gamma$  at time  $t$  equal to  $h\beta_t/2$ , we have

$$X_{t+h} = X_t + h \frac{\beta_t}{2} \nabla \log p(X_t) + \sqrt{h\beta_t} \varepsilon_t.$$

As previously, this is an instance of the Euler-Maruyama scheme for the SDE

$$dX_t = \frac{\beta_t}{2} \nabla \log p(X_t) dt + \sqrt{\beta_t} dW_t, \quad (19)$$

which is nothing but a continuous-time Langevin dynamics! Actually, the VP-SDE, which we also said to be called Langevin equation, corresponds to the continuous Langevin dynamics for a distribution  $p$ , which score function is equal to  $-x$ . This is a standard normal distribution, which automatically explains (in not the most obvious way) why the discrete VP process converges to  $\mathcal{N}(0, I)$ .

### 2.2.3 Continuity equation

ODEs and SDEs describe evolution of the particle's position over time. One of the fundamental questions that one could ask next is how can we describe evolution of the distribution of the particle? In this section, we will answer this question.

Since SDEs only describe change in the position of the particle, one also needs to define the initial point  $X_0$ . Here we assume that it is generated from a distribution  $p_0$  and  $X_0$  is independent with all the noise that comes from the Wiener process. We consider SDEs with scalar diffusion coefficient  $g(t)$  and define the corresponding system

$$\begin{cases} dX_t = f(X_t, t) dt + g(t) dW_t, \\ X_0 \sim p_0. \end{cases} \quad (20)$$

Our goal is to describe how the distribution  $p_t(x) := p_{X_t}(x)$  of the particle changes in time, which almost always means calculating derivative:

$$\frac{\partial}{\partial t} p_t(x) = ?$$

Since we work with SDEs through the discretization, the most convenient way to analyze this time-derivative is through the limit:

$$\frac{\partial}{\partial t} p_t(x) = \lim_{h \rightarrow 0} \frac{1}{h} (p_{t+h}(x) - p_t(x)).$$

Thus, the problem reduces to analyze how density changes after a small time interval. Remember that the variables are connected through the Euler-Maruyama scheme:

$$X_{t+h} = X_t + hf(X_t, t) + \sqrt{h}g(t)\varepsilon_t,$$

where, as always,  $\varepsilon_t \sim \mathcal{N}(0, I)$  is independent with  $X_t$ . This connection is, in fact, pretty simple: the first ODE part  $X_t + hf(X_t, t)$  is nothing but a differentiable function of  $X_t$ , for which (if it is bijective and the inverse is differentiable) there is a change of variables formula. The second part consists in adding an independent noise, which changes density by a convolution.

The first part consists of applying the function  $\varphi(x) = x + hf(x, t)$  to  $X_t$ . Remember that this is an approximation of the function, which solves ODE forward from time  $t$  to time  $t + h$ . ODE theory says that under regularity conditions this function is bijective (uniqueness of the solution) and differentiable in both directions. This allows us to apply the change of variables formula:

$$p_{\varphi(X_t)}(y) = p_{X_t}(\varphi^{-1}(y)) \left| \det \frac{\partial \varphi^{-1}}{\partial y} \right|.$$

The same formula applied in the opposite direction gives

$$p_{X_t}(x) = p_{\varphi(X_t)}(\varphi(x)) \left| \det \frac{\partial \varphi}{\partial x} \right|,$$

which will be more convenient to analyze. We substitute  $\varphi$  and obtain

$$p_{X_t}(x) = p_{\varphi(X_t)}(x + hf(x, t)) \left| \det \left( I + h \frac{\partial}{\partial x} f(x, t) \right) \right|.$$

Let's simplify the det term first. We are interested in all terms that are not  $\bar{o}(h)$ . In the determinant formula for a matrix  $A$  one sums terms of the form  $\prod_{i=1}^d a_{i\sigma_i}$  over all permutations  $\sigma$ . If one of the elements is off the diagonal, it has a pair, which gives  $h^2$  in product and is  $\bar{o}(h)$ . This means that the only thing remained is the diagonal term and the whole det expression is equal

$$\prod_{i=1}^d \left( 1 + h \frac{\partial}{\partial x_i} f_i(x, t) \right) + \bar{o}(h)$$

We see that the only terms that are not  $\bar{o}(h)$  after opening brackets are products that contain 0 or 1  $h$ . This gives

$$1 + h \sum_{i=1}^d \frac{\partial}{\partial x_i} f_i(x, t) + \bar{o}(h) = 1 + h \operatorname{div} f(x, t) + \bar{o}(h),$$

where  $\operatorname{div}$  is the divergence operator, which takes a function  $\psi$  and returns  $\sum_{i=1}^d \partial_{x_i} \psi_i(x) = \operatorname{Tr}(\partial_x \psi(x))$ . Taking small enough  $h$ , one can obtain a positive value, so taking the absolute value is not necessary.

The density part is even easier to deal with: use the Taylor expansion:

$$p_{\varphi(X_t)}(x + hf(x, t)) = p_{\varphi(X_t)}(x) + h \left\langle \nabla p_{\varphi(X_t)}(x), f(x, t) \right\rangle + \bar{o}(h).$$

Combining the two parts, we obtain

$$\begin{aligned} & \left( p_{\varphi(X_t)}(x) + h \left\langle \nabla p_{\varphi(X_t)}(x), f(x, t) \right\rangle + \bar{o}(h) \right) \cdot \left( 1 + h \operatorname{div} f(x, t) + \bar{o}(h) \right) = \\ & = p_{\varphi(X_t)}(x) + h \cdot \left( \left\langle \nabla p_{\varphi(X_t)}(x), f(x, t) \right\rangle + p_{\varphi(X_t)}(x) \operatorname{div} f(x, t) \right) + \bar{o}(h). \end{aligned}$$

Expressing scalar product and divergence as sums, we get

$$p_{\varphi(X_t)}(x) + h \sum_{i=1}^d \left( f_i(x, t) \frac{\partial}{\partial x_i} p_{\varphi(X_t)}(x) + p_{\varphi(X_t)}(x) \frac{\partial}{\partial x_i} f_i(x, t) \right) + \bar{o}(h).$$

Expression under brackets is nothing but a derivative of a product, which gives

$$p_{\varphi(X_t)}(x) + h \sum_{i=1}^d \frac{\partial}{\partial x_i} (f_i(x, t) p_{\varphi(X_t)}(x)) + \bar{o}(h).$$

The sum is the divergence of the product  $f(x, t)p_{\varphi(X_t)}(x)$ , which finally gives

$$p_{X_t}(x) = p_{\varphi(X_t)}(x) + h \operatorname{div} \left( f(x, t)p_{\varphi(X_t)}(x) \right) + \bar{o}(h). \quad (21)$$

Remember that  $\varphi(X_t) = X_t + hf(X_t, t)$ . If we work with ODE, then  $\varphi(X_t) = X_{t+h}$  and

$$p_t(x) = p_{t+h}(x) + h \operatorname{div} \left( f(x, t)p_{t+h}(x) \right) + \bar{o}(h),$$

which is equivalent to

$$\frac{p_{t+h}(x) - p_t(x)}{h} = -\operatorname{div} \left( f(x, t)p_{t+h}(x) \right) + \bar{o}(1).$$

Taking the limit, we obtain

$$\boxed{\frac{\partial}{\partial t} p_t(x) = -\operatorname{div} \left( f(x, t)p_t(x) \right)} \quad (22)$$

which is called the *continuity equation* and describes evolution of a particle that moves under ODE.

#### 2.2.4 Fokker-Planck equation

The continuity equation is a very important object which can be investigated separately, but we also need its extension on the SDE case. Previously, we expressed  $X_{t+h}$  through the Euler-Maruyama scheme:

$$X_{t+h} = \varphi(X_t) + Y_t,$$

where  $Y_t = \sqrt{h}g(t)\varepsilon_t$ , and found density of  $\varphi(X_t)$  in the Equation 21. Since  $\varphi(X_t)$  and  $Y_t$  are independent, density of  $X_{t+h}$  is just a convolution

$$p_{X_{t+h}}(x) = \int p_{\varphi(X_t)}(x-y)p_{Y_t}(y)dy = \mathbb{E} p_{\varphi(X_t)}(x - Y_t) = \mathbb{E} p_{\varphi(X_t)}(x - \sqrt{h}g(t)\varepsilon_t)$$

As always, use the Taylor expansion (up to the second term, since  $\sqrt{h}^2 = h$ ):

$$\begin{aligned} \mathbb{E} \left( p_{\varphi(X_t)}(x) - \left\langle \nabla p_{\varphi(X_t)}(x), \sqrt{h}g(t)\varepsilon_t \right\rangle + \frac{1}{2} \left\langle \nabla^2 p_{\varphi(X_t)}(x) \sqrt{h}g(t)\varepsilon_t, \sqrt{h}g(t)\varepsilon_t \right\rangle \right) + \\ + \mathbb{E} \bar{o} \left( hg^2(t) \|\varepsilon_t\|^2 \right). \end{aligned}$$

This expression is large but convenient to work with:

1. The first summand is deterministic, so  $\mathbb{E} p_{\varphi(X_t)}(x) = p_{\varphi(X_t)}(x)$ ;
2. The second is a linear function of a zero-mean variable, which also has zero mean:

$$\mathbb{E} \left\langle \nabla p_{\varphi(X_t)}(x), \sqrt{h}g(t)\varepsilon_t \right\rangle = \left\langle \nabla p_{\varphi(X_t)}(x), \sqrt{h}g(t)\mathbb{E}\varepsilon_t \right\rangle = 0$$

3. The last term under expectation is  $\bar{o}(h)$ . By interchanging small- $\bar{o}$  and expectation (which is a very rude operation), we obtain  $\bar{o}(h)$ .

Now, we have a much smaller expression

$$\begin{aligned} p_{\varphi(X_t)}(x) + \mathbb{E} \left( \frac{1}{2} \left\langle \nabla^2 p_{\varphi(X_t)}(x) \sqrt{h}g(t)\varepsilon_t, \sqrt{h}g(t)\varepsilon_t \right\rangle \right) + \bar{o}(h) = \\ = p_{\varphi(X_t)}(x) + \frac{hg^2(t)}{2} \mathbb{E} \left\langle \nabla^2 p_{\varphi(X_t)}(x) \varepsilon_t, \varepsilon_t \right\rangle + \bar{o}(h). \end{aligned}$$

The term under expectation here is the famous Hutchinson's trace estimator [16]. Rewrite it using a trace cyclic property:

$$\frac{hg^2(t)}{2} \mathbb{E} \left( \varepsilon_t^\top \nabla^2 p_{\varphi(X_t)}(x) \varepsilon_t \right) = \frac{hg^2(t)}{2} \mathbb{E} \operatorname{Tr} \left( \nabla^2 p_{\varphi(X_t)}(x) \varepsilon_t \varepsilon_t^\top \right).$$

Trace is a linear function, which we can interchange with expectation and obtain

$$\frac{hg^2(t)}{2} \text{Tr} \left( \nabla^2 p_{\varphi(X_t)}(x) \mathbb{E} \varepsilon_t \varepsilon_t^\top \right).$$

Finally, since  $\mathbb{E} \varepsilon_t \varepsilon_t^\top = \text{Var} \varepsilon_t + (\mathbb{E} \varepsilon_t)(\mathbb{E} \varepsilon_t)^\top$  and  $\varepsilon_t \sim \mathcal{N}(0, I)$ , it is equal to  $\text{Var} \varepsilon_t = I$ . We obtained

$$\begin{aligned} p_{X_{t+h}}(x) &= p_{\varphi(X_t)}(x) + \frac{hg^2(t)}{2} \text{Tr} \left( \nabla^2 p_{\varphi(X_t)}(x) \right) + \bar{o}(h) = \\ &= p_{\varphi(X_t)}(x) + \frac{hg^2(t)}{2} \Delta p_{\varphi(X_t)}(x) + \bar{o}(h), \end{aligned}$$

where  $\Delta \psi(x) = \sum_{i=1}^d \partial_{x_i x_i}^2 \psi(x)$  is the Laplace operator. Rewriting density of  $\varphi(X_t)$  as we expressed in the Equation 21, we obtain

$$p_{X_{t+h}}(x) = p_{X_t}(x) - h \text{div} \left( f(x, t) p_{\varphi(X_t)}(x) \right) + \frac{hg^2(t)}{2} \Delta p_{\varphi(X_t)}(x) + \bar{o}(h),$$

which is equivalent to

$$\frac{p_{t+h}(x) - p_t(x)}{h} = -\text{div} \left( f(x, t) p_{\varphi(X_t)}(x) \right) + \frac{g^2(t)}{2} \Delta p_{\varphi(X_t)}(x) + \bar{o}(1).$$

Since  $\varphi(X_t) = X_t + hf(X_t, t) \xrightarrow{h \rightarrow 0} X_t$ , we take the limit  $h \rightarrow 0$  and obtain

$$\boxed{\frac{\partial}{\partial t} p_t(x) = -\text{div} \left( f(x, t) p_t(x) \right) + \frac{g^2(t)}{2} \Delta p_t(x)} \quad (23)$$

which is called the *Fokker-Planck equation* and describes how distribution of the particle, driven by SDE, evolves in time.

### 2.2.5 Uniqueness and physical interpretation

Continuity (CE) and Fokker-Planck (FPE) equations are fundamental results which demonstrate some properties of the distribution dynamics: it should coincide with Equation 22 in case of ODE and Equation 23 in case of SDE. However, it is not clear from the first glance, whether the distribution dynamics that solves FPE(CE) is unique and under which initial conditions. Can we say, for example, that if the density dynamics satisfies FPE with velocity field  $f$  and diffusion coefficient  $g$ , then this dynamics is generated by SDE with the corresponding parameters?

The answer is positive: under some conditions, the Cauchy problem for FPE

$$\begin{cases} \frac{\partial}{\partial t} p_t(x) = -\text{div} \left( f(x, t) p_t(x) \right) + \frac{g^2(t)}{2} \Delta p_t(x); \\ p_0(x) = q(x) \end{cases} \quad (24)$$

has a unique solution. This result and many others can be found, for example, in [3] (continuity equation) and [6, 5] (Fokker-Planck equation).

This has a very useful practical consequence: saying that  $p_t(x)$  satisfies FPE with initial condition, as in Equation 24, is equivalent to say that  $p_t(x)$  is generated by

$$\begin{cases} dX_t = f(X_t, t)dt + g(t)dW_t; \\ X_0 \sim q. \end{cases} \quad (25)$$

This allows to analyze properties of trajectories by studying the corresponding FPE and analyze solution of the FPE by studying the trajectories, which we will use multiple times.

Last, but not the least, continuity and Fokker-Planck equations have a meaningful physical interpretation. This is easier to demonstrate in 1-dimensional case, where FPE rewrites as

$$\frac{\partial}{\partial t} p_t(x) = -\frac{\partial}{\partial x} \left( f(x, t) p_t(x) \right) + \frac{g^2(t)}{2} \frac{\partial^2}{\partial x^2} p_t(x) = -\frac{\partial}{\partial x} \left( f(x, t) p_t(x) - \frac{g^2(t)}{2} \frac{\partial}{\partial x} p_t(x) \right).$$

Next derivations are adapted from the PDE course by S. Shaposhnikov and T. Krasovitsky and can be found here. Let us consider a function  $u_t(x)$ , which defines concentration of some substance at time  $t$  and assume that  $F(x, t)$  is a so-called flow, which defines the amount of substance that flowed in the point  $x_0$  from time  $t_0$  to  $t_0 + \Delta t$  by the formula

$$\int_{t_0}^{t_0+\Delta t} F(x, t) dt.$$

By definition, the amount of substance contained in a segment  $[x_0, x_0 + \Delta x]$  is equal to

$$\int_{x_0}^{x_0+\Delta x} u_t(x) dx.$$

On the one hand, difference between the amount of substance at time  $t_0$  and  $t_0 + \Delta t$  is equal to

$$\int_{x_0}^{x_0+\Delta x} u_{t_0+\Delta t}(x) dx - \int_{x_0}^{x_0+\Delta x} u_{t_0}(x) dx.$$

On the other hand, it is equal to the difference between the amount that flowed in  $x_0$  and flowed out of  $x_0 + \Delta x$ :

$$\int_{t_0}^{t_0+\Delta t} F(x_0, t) dt - \int_{t_0}^{t_0+\Delta t} F(x_0 + \Delta x, t) dt,$$

which means that

$$\int_{x_0}^{x_0+\Delta x} u_{t_0+\Delta t}(x) dx - \int_{x_0}^{x_0+\Delta x} u_{t_0}(x) dx = \int_{t_0}^{t_0+\Delta t} F(x_0, t) dt - \int_{t_0}^{t_0+\Delta t} F(x_0 + \Delta x, t) dt.$$

By using the Newton-Leibnitz theorem inside integrals, we obtain

$$\int_{x_0}^{x_0+\Delta x} \int_{t_0}^{t_0+\Delta t} \frac{\partial}{\partial t} u_t(x) dt dx = - \int_{t_0}^{t_0+\Delta t} \int_{x_0}^{x_0+\Delta x} \frac{\partial}{\partial x} F(x, t) dx dt.$$

Dividing by  $\Delta x \Delta t$  and taking  $\Delta x \rightarrow 0, \Delta t \rightarrow 0$ , we obtain that

$$\frac{\partial}{\partial t} u_t(x) = - \frac{\partial}{\partial x} F(x, t)$$

holds for all pairs  $(x, t)$ . This family of equations is called *conservation laws* and describes how the amount of substance changes in time, when the substance flow is equal to  $F$ . Trivially, continuity and Fokker-Planck equations are instances of the conservation laws with

$$F(x, t) = f(x, t) p_t(x)$$

for continuity equation and

$$F(x, t) = f(x, t) p_t(x) - \frac{g^2(t)}{2} \frac{\partial}{\partial x} p_t(x)$$

for FPE. In both cases we treat probability density  $p_t(x)$  as concentration of the substance. In the first case, flow direction coincides with the velocity field and the flow value is proportional to the concentration in the point. In the second case, the direction is modified by subtracting gradient of the density (with coefficient), which means that the substance should flow in the direction, where concentration is low and, thus, compensate it. The obtained flow represents a combination of deterministic force and diffusion, which tends to spread out.

### 2.3 Diffusion models based on SDEs

Now that we are familiar with SDEs and Fokker-Planck equations, we are ready to construct diffusion models in continuous time, as was done in [25]. Let  $X_t$  be the forward noising process, defined by the system

$$\begin{cases} dX_t = f(X_t, t)dt + g(t)dW_t; \\ X_0 \sim p_{data}, \end{cases}$$

such that  $p_1$  is very close to some simple distribution as, for example,  $\mathcal{N}(0, I)$ . The most common choice is the VP-SDE, defined in the Equation 18.

Our goal is to somehow construct the reverse process, which translates a known simple distribution  $p_1$  into  $p_{data}$ . The first observation is that if we worked with ODEs instead of SDEs, this would be trivial. If  $X_t$  is generated by an ODE  $dX_t = f(X_t, t)dt$  and  $Y_t = X_{1-t}$  is the corresponding reverse process, then

$$\frac{d}{dt}Y_t = \frac{d}{dt}X_{1-t} = -\frac{\partial}{\partial t}X_{1-t} = -f(X_{1-t}, 1-t) = -f(Y_t, 1-t),$$

which means that the reverse process is also generated by ODE with the velocity field looking in the opposite direction.

This leads to a natural idea: turn SDE into an ODE, reverse it and turn into SDE back! But what does it mean to turn SDE into ODE, in what sense? In Section 2.2.5 we said that distribution dynamics is generated by an SDE if and only if this dynamics solves the corresponding FPE/CE. Since in our case the forward process is represented by an SDE, its dynamics solves the corresponding FPE

$$\begin{cases} \frac{\partial}{\partial t}p_t(x) = -\operatorname{div}\left(f(x, t)p_t(x)\right) + \frac{g^2(t)}{2}\Delta p_t(x); \\ p_0 = p_{data}. \end{cases}$$

If we manage to find such velocity field  $v(x, t)$ , that  $p_t(x)$  will also solve the continuity equation

$$\frac{\partial}{\partial t}p_t(x) = -\operatorname{div}\left(v(x, t)p_t(x)\right),$$

then ODE  $dZ_t = v(Z_t, t)dt$  with initial condition  $Z_0 \sim p_{data}$  will generate the same distribution dynamics. Constructing such  $v$  is not a big deal: we should move the second summand in FPE inside the divergence. Note that  $\Delta p_t(x) = \sum_{i=1}^d \partial_{x_i}^2 p_t(x) = \sum_{i=1}^d \partial_{x_i}(\partial_{x_i} p_t(x)) = \operatorname{div} \nabla p_t(x)$ . Then, the FPE reads

$$\frac{\partial}{\partial t}p_t(x) = -\operatorname{div}\left(f(x, t)p_t(x)\right) + \frac{g^2(t)}{2}\operatorname{div} \nabla p_t(x) = -\operatorname{div}\left(f(x, t)p_t(x) - \frac{g^2(t)}{2}\nabla p_t(x)\right).$$

Now we use the beloved log-derivative trick and obtain

$$-\operatorname{div}\left(f(x, t)p_t(x) - \frac{g^2(t)}{2}\nabla \log p_t(x) p_t(x)\right) = -\operatorname{div}\left(\left(f(x, t) - \frac{g^2(t)}{2}\nabla \log p_t(x)\right)p_t(x)\right),$$

which means that  $p_t(x)$ , generated by the noising process, solves the continuity equation with velocity field

$$v(x, t) = f(x, t) - \frac{g^2(t)}{2}\nabla \log p_t(x)$$

and, thus, can be generated by solving

$$\begin{cases} dY_t = v(Y_t, t)dt; \\ Y_0 \sim p_{data}. \end{cases}$$

The more interesting for us result is that the backward dynamics  $q_t(x) = p_{1-t}(x)$  can be generated by the reverse process  $Y_t^{(b)} = Y_{1-t}$ , which, as we discussed earlier, corresponds to the ODE

$$\begin{cases} dY_t^{(b)} = -v(Y_t^{(b)}, 1-t)dt; \\ Y_0^{(b)} \sim q_0. \end{cases}$$

Remember that  $q_0 = p_1$  is a very simple distribution, e.g. close to  $\mathcal{N}(0, I)$ . Thus, we already achieved a way to generate the backward dynamics, however, did it through ODE instead of SDE. And the backward ODE has the score function inside the velocity field, which clearly demonstrates its importance from another point of view, different from Langevin dynamics.

The original goal, however, was in constructing the backward SDE and not ODE. Let's return the favor. We know that  $q_t$  solves the continuity equation

$$\frac{\partial}{\partial t} q_t(x) = -\operatorname{div}\left(-v(x, 1-t)q_t(x)\right)$$

Now the only thing to do is to extract the Laplace operator with the coefficient  $g^2(1-t)/2$  from the brackets, which will correspond to the original diffusion coefficient, but with the reversed time. By adding and subtracting it, we obtain

$$\begin{aligned} & -\operatorname{div}\left(-v(x, 1-t)q_t(x)\right) - \frac{g^2(1-t)}{2}\operatorname{div}\nabla q_t(x) + \frac{g^2(1-t)}{2}\Delta q_t(x) = \\ & = -\operatorname{div}\left(-v(x, 1-t)q_t(x) + \frac{g^2(1-t)}{2}\nabla q_t(x)\right) + \frac{g^2(1-t)}{2}\Delta q_t(x) = \\ & = -\operatorname{div}\left(\left(-v(x, 1-t) + \frac{g^2(1-t)}{2}\nabla \log q_t(x)\right)q_t(x)\right) + \frac{g^2(1-t)}{2}\Delta q_t(x). \end{aligned}$$

Now we simplify the velocity field:

$$\begin{aligned} & -v(x, 1-t) + \frac{g^2(1-t)}{2}\nabla \log q_t(x) = \\ & = -f(x, 1-t) + \frac{g^2(1-t)}{2}\nabla \log p_{1-t}(x) + \frac{g^2(1-t)}{2}\nabla \log q_t(x) = \\ & = -f(x, 1-t) + g^2(1-t)\nabla \log p_{1-t}(x). \end{aligned}$$

Thus, we showed that  $q_t(x) = p_{1-t}(x)$  solves the Fokker-Planck equation with diffusion coefficient  $g^2(1-t)$  and velocity field  $-f(x, 1-t) + g^2(1-t)\nabla \log p_{1-t}(x)$  and can be generated by solving

$$\begin{cases} dX_t^{(b)} = \left(-f\left(X_t^{(b)}, 1-t\right) + g^2(1-t)\nabla \log p_{1-t}(X_t^{(b)})\right) dt + g(1-t)dW_t; \\ X_0^{(b)} \sim p_1. \end{cases} \quad (26)$$

We proved that  $(X_t^{(b)}, t \in [0, 1])$  is the reverse process for the  $(X_t, t \in [0, 1])$  only in terms of marginal distributions at all time steps:

$$p_{X_{1-t}} = p_{X_t^{(b)}},$$

since FPE does not account for any joint distributions. However, there is a much more powerful result that can be found in [4], which states that the whole processes have the same distribution: for all measurable  $B \subset \mathcal{C}[0, 1]$

$$\mathbb{P}\left((X_{1-t}, t \in [0, 1]) \in B\right) = \mathbb{P}\left((X_t^{(b)}, t \in [0, 1]) \in B\right).$$

### 2.3.1 Practical usage

Going back into the real world, we can now sample from a reverse process by first sampling  $X_0^{(b)} \sim p_1$  and solving the backward SDE 26 with Euler-Maruyama scheme:

$$X_{t+h}^{(b)} \approx X_t^{(b)} - hf(X_t^{(b)}, 1-t) + hg^2(1-t)\nabla \log p_{1-t}(X_t^{(b)}) + \sqrt{h}g(1-t)\varepsilon_t.$$

Or one can write it in a more common backward way: sample  $X_1 \sim p_1$  and iterate

$$X_{t-h} \approx X_t - hf(X_t, t) + hg^2(t)\nabla \log p_t(X_t) + \sqrt{h}g(t)\varepsilon_t.$$

Of course, we do not know the score-function  $\nabla \log p_t(x)$ , that is why we learn it with the score matching:

$$\int_0^1 \mathbb{E} \|s_\theta(X_t, t) - \nabla \log p_{t|0}(X_t|X_0)\|^2 dt \rightarrow \min_\theta, \quad (27)$$

where  $X_0 \sim p_{data}$  and  $X_t$  is the forward process. As previously, we replace functional with a Monte Carlo estimate and calculate gradient as

$$\nabla_\theta \|s_\theta(X_t, t) - \nabla \log p_{t|0}(X_t|X_0)\|^2,$$

where  $t \sim \mathcal{U}[0, 1]$ . To efficiently optimize the functional, one needs to be able to calculate score function of  $p_{t|0}$  and sample from it. One can take, for example, the VP-SDE, defined in the Equation 18, which conditional distribution  $p_{t|0}$  is normal (can be obtained by using the Euler-Maruyama scheme and taking the corresponding limit). The sampling scheme then will look as generating  $X_1 \sim \mathcal{N}(0, I)$  and iterating

$$X_{t-h} = X_t \left(1 + h \frac{\beta_t}{2}\right) + h\beta_t \cdot s_\theta(X_t, t) + \sqrt{h\beta_t} \varepsilon_t.$$

Compare it with the Langevin dynamics, which is used to sample from NCSN (taking the step size  $h\beta_t$ ):

$$X_t^{(m+1)} = X_t^{(m)} + h\beta_t \cdot s_\theta(X_t^{(m)}, t) + \sqrt{2h\beta_t} \varepsilon_t^{(m)}.$$

Ideologically, both schemes represent the same idea: one should perform a random walk balancing between direction of the fastest growth and chaotic Brownian motion. The two differences between steps of the schemes are noise term coefficient (which is two times bigger in Langevin) and the additional term  $-hf(X_t, t)$  (or  $h\beta_t X_t/2$  for VP-SDE) in Euler scheme for backward diffusion. These slight changes, however, allow to perform just one step per noise level, which reduces the overall number of hyperparameters and seems more natural, since distributions, adjacent in time, are very similar.

### 2.3.2 Calculating likelihood

While constructing the reverse diffusion process, we first turned the original SDE into ODE with equivalent marginals because ODEs are easy to reverse. This, however, is a very powerful technique itself: it allows to treat an SDE diffusion model as a continuous normalizing flow [7, 12], function, which is a diffeomorphism (bijective differentiable mapping with differentiable inverse, e.g. change of variables) between a simple distribution  $p_1$  such as  $\mathcal{N}(0, I)$  and  $p_{data}$ . This class of functions is very convenient to work with, since it allows to calculate density of the transformed variable by calculating the original density and the corresponding Jacobian:

$$p_{\varphi(X)}(y) = p_X(\varphi^{-1}(y)) \left| \det \frac{\partial \varphi^{-1}}{\partial y} \right|.$$

In this case, however,  $\varphi$  is quite a complex function that takes the initial value and solves the corresponding ODE. That's why we will not use this formula directly and take a look from the other side.

Assume we work with an ODE-based model

$$\begin{cases} dY_t = v(Y_t, t)dt; \\ Y_0 \sim p_0, \end{cases}$$

such that  $Y_1$  has a simple tractable distribution  $p_1$ . This can be, for example, ODE version of the forward diffusion process. Our aim is to construct an algorithm of calculating likelihood  $\log p_0(Y_0)$  of any given image  $Y_0$ . Solving ODE forward in time allows us to obtain the sample  $Y_1$ , for which calculating likelihood  $\log p_1(Y_1)$  is easy, since  $p_1$  is easy. The question is how these two values are connected. Remember that evolution of the density is described by the continuity equation

$$\frac{\partial}{\partial t} p_t(x) = -\text{div}(v(x, t)p_t(x)).$$



One can rewrite it with respect to the logarithm:

$$\begin{aligned}\frac{\partial}{\partial t} \log p_t(x) &= \frac{\frac{\partial}{\partial t} p_t(x)}{p_t(x)} = \frac{-\operatorname{div}(v(x, t)p_t(x))}{p_t(x)} = \frac{-\sum_{i=1}^d \left( p_t(x) \frac{\partial}{\partial x_i} v_i(x, t) + v_i(x, t) \frac{\partial}{\partial x_i} p_t(x) \right)}{p_t(x)} = \\ &= -\frac{1}{p_t(x)} \left( p_t(x) \operatorname{div} v(x, t) + \langle v(x, t), \nabla p_t(x) \rangle \right) = -\langle v(x, t), \nabla \log p_t(x) \rangle - \operatorname{div} v(x, t).\end{aligned}$$

This equation, being an equivalent for the continuity equation, describes evolution of the log-density in an arbitrary point. We would like, however, to describe evolution of the log-density on the trajectory  $Y_t$ . As always, describing evolution means finding derivative:

$$\frac{d}{dt} \log p_t(Y_t) = \frac{\partial}{\partial t} \log p_t(x) \Big|_{x=Y_t} + \left\langle \nabla \log p_t(Y_t), \frac{\partial}{\partial t} Y_t \right\rangle.$$

Here  $d/dt$  denotes the full derivative,  $\partial/\partial t$  the partial, and we used the law of differentiating the composition of functions. Using the equation for logarithm and derivative of  $Y_t$ , we obtain

$$\frac{d}{dt} \log p_t(Y_t) = -\langle v(Y_t, t), \nabla \log p_t(Y_t) \rangle - \operatorname{div} v(Y_t, t) + \langle \nabla \log p_t(Y_t), v(Y_t, t) \rangle = -\operatorname{div} v(Y_t, t),$$

which is exactly what we were looking for! This substitution of the trajectory inside the partial differential equation is called *characteristic method* and is a key idea for solving equations of such type. Now, we can extend the original ODE for the particle with ODE for the likelihood and obtain the system

$$\begin{cases} \frac{d}{dt} Y_t = v(Y_t, t); \\ \frac{d}{dt} L_t := \frac{d}{dt} \log p_t(Y_t) = -\operatorname{div} v(Y_t, t), \end{cases}$$

which can be approximately solved, for example, by the Euler scheme

$$\begin{cases} Y_{t+h} \approx Y_t + h v(Y_t, t); \\ L_{t+h} \approx L_t - h \operatorname{div} v(Y_t, t) \end{cases}$$

and allows us to calculate likelihood of the particle along with its position. But wait a second: our goal is to calculate  $L_0$  knowing  $L_1$ , but here we iterate  $L_t$  forward in time, assuming we know  $L_0$ . The key feature of the joint ODE is that the right-hand side, corresponding to  $L_t$ , does not depend on  $L_t$ . This means that one can calculate the difference between  $L_0$  and  $L_1$  without knowing  $L_0$ . Here, for example, one can iterate the full Euler scheme and obtain

$$L_0 \approx L_h + h \operatorname{div} v(Y_0, 0) \approx L_{2h} + h \operatorname{div} v(Y_h, h) + h \operatorname{div} v(Y_0, 0) \approx \dots \approx L_1 + h \sum_t \operatorname{div} v(Y_t, t),$$

where  $L_1 = \log p_1(Y_1)$  is easy to calculate, since  $p_1$  is simple and  $Y_1$  is obtained by solving the ODE forward. The algorithm, thus, consists in solving  $dY_t = v(Y_t, t)dt$  forward along with calculating sum of the divergences  $\operatorname{div} v(Y_t, t)$  and likelihood  $L_1$  of the final sample.

This algorithm, however, has a serious practical limitation: one should efficiently calculate  $\operatorname{div} v(x, t)$ , which is trace of the Jacobian of the neural network, which has  $(H \times W \times 3)^2$  elements and is intractable for most practical cases. This means that one should use algorithms that allow to calculate trace without computing the matrix. One of the most suitable choices here is the Hutchinson's trace estimator [16], which we already faced while deriving Fokker-Planck equation. Let  $\varepsilon \in \mathbb{R}^d$  be a random vector with  $\mathbb{E} \varepsilon = 0$  and  $\operatorname{Var} \varepsilon = I$ . Then for a  $d \times d$  matrix  $A$

$$\mathbb{E} \langle \varepsilon, A\varepsilon \rangle = \mathbb{E} \varepsilon^\top A \varepsilon = \mathbb{E} \operatorname{Tr} (\varepsilon^\top A \varepsilon) = \mathbb{E} \operatorname{Tr} (A \varepsilon \varepsilon^\top) = \operatorname{Tr} (A \mathbb{E} \varepsilon \varepsilon^\top) = \operatorname{Tr} (A \mathbb{E} \varepsilon \varepsilon^\top).$$

Here we used the trace cyclic property (second equality) and linearity of the expectation, when interchanging it with trace and matrix multiplication (third and fourth equalities). Since  $\operatorname{Var} \varepsilon = \mathbb{E} \varepsilon \varepsilon^\top - (\mathbb{E} \varepsilon)(\mathbb{E} \varepsilon)^\top$  and  $\mathbb{E} \varepsilon = 0$ , the term  $\mathbb{E} \varepsilon \varepsilon^\top$  is just the identity matrix and we obtain

$$\operatorname{Tr} A = \mathbb{E} \langle \varepsilon, A\varepsilon \rangle,$$

which means that  $\langle \varepsilon, A\varepsilon \rangle$  is an unbiased estimate of the trace of the matrix  $A$ . In particular, one can control the estimation error by sampling the necessary amount  $m$  of independent  $\varepsilon_i$  from a suitable distribution (such as  $\mathcal{N}(0, I)$  or a vector of independent random signs) and estimating trace as  $\sum_{i=1}^m \langle \varepsilon_i, A\varepsilon_i \rangle / m$ .

Going back to divergences, we modify the described algorithm with estimating

$$\operatorname{div} v(x, t) \approx \frac{1}{m} \sum_{i=1}^m \left\langle \varepsilon_i, \frac{\partial}{\partial x} v(x, t) \varepsilon_i \right\rangle.$$

This estimator reduces the problem of calculating trace of the jacobian to calculating jacobian-vector product, which can be relatively efficiently implemented by backpropagation.

## 2.4 Conditional diffusion models

### 2.4.1 Conditional training

At this moment, we constructed more or less everything that is needed for training and generating samples with diffusion models. A natural next step would be to generalize diffusion models for conditional tasks such as generating samples according to some class label or textual prompt.

Formally, let  $(X_0, Y)$  be a pair from the dataset, where  $Y$  is a variable of arbitrary nature, and  $X_t$  is, as always, a forward process (discrete or continuous-time), where all the noise is independent from  $Y$ . For simplicity, we will call  $Y$  label, regardless of its meaning. The problem of conditional generation consists in generating samples from  $p_{X_0|Y}$ . Without noising processes, one could do this by learning the corresponding score function  $\nabla \log p_{X_0|Y}(x_0|y)$ .

In the Section 2.1.1, however, we discussed why it is a bad idea and motivated usage of the noising process. Since both sampling algorithms, backward SDE and Langevin dynamics, depend only on the loss function (this is, in fact, incorrect, but can be considered as a heuristic for now), one can just replace score functions of the noising process  $\nabla \log p_{X_t}(x_t)$  with the corresponding conditional score functions  $\nabla \log p_{X_t|Y}(x_t|y)$  and learn them. Analogous to derivations in the Section 2.1.2, these  $Y$ -conditional score functions can be represented as conditional expectation of the score function, conditioned on both  $Y$  and  $X_0 = X$ :

$$\nabla \log p_{X_t|Y}(x_t|y) = \mathbb{E} [\nabla \log p_{X_t|X_0,Y}(X_t|X_0, Y) \mid X_t = x_t, Y = y]$$

and, thus, can be trained by solving

$$\mathbb{E} \|s_\theta(X_t, t|Y) - \nabla \log p_{X_t|X_0,Y}(X_t|X_0, Y)\|^2 \rightarrow \min_\theta.$$

The good thing is that the conditional distribution  $p_{X_t|X_0,Y}$  can be simplified into  $p_{X_t|X_0}$ , since, by construction,  $X_t$  is completely determined by  $X_0$  and the added noise, which is independent from  $Y$  by construction. Thus, the conditional score functions for all time steps can be jointly trained as (with sum instead of integral for discrete-time processes):

$$\int_0^1 \mathbb{E} \|s_\theta(X_t, t|Y) - \nabla \log p_{t|0}(X_t|X_0)\|^2 dt \rightarrow \min_\theta.$$

The resulting loss is a very simple extension of the original denoising score matching from the Equation 27. The only difference is that the model now takes label  $Y$  as the additional input and tries to denoise the image using the additional information in  $Y$ . But what if we have a trained unconditional diffusion model and want to perform a conditional generation?

### 2.4.2 Classifier guidance

The next technique is called *classifier guidance*. It was introduced in [25] and used to obtain state-of-the-art in [9]. The idea is simple yet elegant: rewrite the conditional score using the Bayes' formula:

$$\begin{aligned} \nabla \log p_{X_t|Y}(x_t|y) &= \nabla (\log p_{X_t}(x_t) + \log p_{Y|X_t}(y|x_t) - \log p_Y(y)) = \\ &= \nabla \log p_{X_t}(x_t) + \nabla \log p_{Y|X_t}(y|x_t), \end{aligned}$$

since  $\log p_Y(y)$  does not depend on  $x_t$ . The first term can be approximated with a trained unconditional model  $s_\theta(x_t, t)$ . The second term is simply score of the distribution of the label corresponding to the noisy image. If  $Y$  is a class label, then this distribution can be approximated by a classifier  $c(y|x_t, t)$ , which, however, should also differ with changing time  $t$  and take noisy image  $x_t$  as an

input. Classifiers of this type are called *noisy classifiers*. It should be noted that the same logic works for arbitrary complex structures  $Y$ , for which there is a model that calculates probability given a noisy sample and time. Having such model enables conditional sampling for unconditional models by approximating

$$\nabla \log p_{X_t|Y}(x_t|y) \approx s_\theta(x_t, t) + \nabla \log c(y|x_t, t).$$

The best practical results, however, are obtained by introducing and varying the additional classifier weight  $\gamma$  in the formula

$$s_\theta(x_t, t) + \gamma \nabla \log c(y|x_t, t),$$

which allows to control rate of the classifier signal. Big values of  $\gamma$  correspond to generation of the most probable sample according to the classifier. Small values of  $\gamma$ , in contrast, lead to almost unconditional samples. This means that controlling  $\gamma$ , one can solve the *diversity vs fidelity* tradeoff and generate diverse picture with high quality or, at least, optimally choose between the two. In practice, moderate values of  $\gamma$  such as 5 – 10 improve generative metrics over  $\gamma = 1$ , which justifies this modified procedure.

The funny thing is that the quality can be improved [9] even further, if an already conditional model is combined with the classifier guidance as

$$s_\theta(x_t, t|y) + \gamma \nabla \log c(y|x_t, t).$$

This means that extracting information from label is valuable in both ways and that the two techniques do not contradict one another. It also allows to control the diversity-fidelity tradeoff for class-conditional models too, which is a valuable outcome.

### 2.4.3 Classifier-free guidance

As we see, classifier guidance, being originally a technique for conditional sampling from unconditional models, became a very useful tool for conditional models because of the possibility to choose  $\gamma > 1$ . The overall procedure, however, becomes sophisticated due to the presence of the noisy classifier. If  $Y$  is a class label and  $c$  is indeed a classifier, then it is accessible, however, if  $Y$  is a text prompt, one needs to train an additional language model, conditioned on noisy samples. On the other hand, it seems that the conditional distribution  $p_{Y|X_t}$  can be extracted from a conditional diffusion model  $p_{X_t|Y}$  without any additional classifiers. This idea leads to the *classifier-free guidance*, introduced in [15]. The first step is to take classifier weight equal to  $1 + \gamma$  (since this is an interesting case for  $\gamma > 0$ ) and write

$$\begin{aligned} & \nabla \log p_{X_t}(x_t) + (1 + \gamma) \nabla \log p_{Y|X_t}(y|x_t) = \\ & = \nabla \log p_{X_t}(x_t) + (1 + \gamma) (\nabla \log p_{X_t|Y}(x_t|y) - \nabla \log p_{X_t}(x_t)) = \\ & = \nabla \log p_{X_t|Y}(x_t|y) + \gamma (\nabla \log p_{X_t|Y}(x_t|y) - \nabla \log p_{X_t}(x_t)). \end{aligned}$$

Thus, we represented the classifier-weighted conditional score as sum of the original conditional score and weighted difference between the conditional and unconditional scores. Given a class-conditional model, one can replace  $\nabla \log p_{X_t|Y}(x_t|y) \approx s_\theta(x_t, t|y)$ . But what to do with unconditional  $\nabla \log p_{X_t}(x_t)$ ? The idea is to say that the unconditional score  $\nabla \log p_{X_t}(x_t)$  is the conditional score, but with the specific label  $\emptyset$ , with which the diffusion model is trained as unconditional:

$$\begin{aligned} & \nabla \log p_{X_t|Y}(x_t|y) + \gamma (\nabla \log p_{X_t|Y}(x_t|y) - \nabla \log p_{X_t}(x_t)) \approx \\ & s_\theta(x_t, t|y) + \gamma (s_\theta(x_t, t|y) - s_\theta(x_t, t|\emptyset)). \end{aligned}$$

The formula itself is very simple and meaningful: the first term corresponds to the ordinary class-conditional generation. The second term has a weight  $\gamma$ , the additional weight of the classifier, which is multiplied by the direction, which distinguishes conditional generation from unconditional.

The training procedure remains the same, but the label  $Y$  is replaced with  $\emptyset$  with some probability, which corresponds to training  $s_\theta(x_t, t|\emptyset)$  on  $\nabla \log p_{X_t}(x_t)$ . This approach is very general and can be applied for arbitrary structures: if  $Y$  is a class label, then  $\emptyset$  is an additional class, corresponding to the absence of the class; in texts,  $\emptyset$  corresponds to the empty prompt, etc. This is a simple, yet very powerful technique, which recently gained a lot of popularity, since it allows to choose between quality and diversity of generated samples without need to train an additional model.

## References

- [1] Michael S Albergo, Nicholas M Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *arXiv preprint arXiv:2303.08797*, 2023.
- [2] Michael S Albergo and Eric Vanden-Eijnden. Building normalizing flows with stochastic interpolants. *arXiv preprint arXiv:2209.15571*, 2022.
- [3] Luigi Ambrosio, Nicola Gigli, and Giuseppe Savaré. *Gradient flows: in metric spaces and in the space of probability measures*. Springer Science & Business Media, 2005.
- [4] Brian DO Anderson. Reverse-time diffusion equation models. *Stochastic Processes and their Applications*, 12(3):313–326, 1982.
- [5] Vladimir I Bogachev, Nicolai V Krylov, Michael Röckner, and Stanislav V Shaposhnikov. *Fokker–Planck–Kolmogorov Equations*, volume 207. American Mathematical Society, 2022.
- [6] Vladimir Igorevich Bogachev, Tikhon Il’ich Krasovitskii, and Stanislav Valer’evich Shaposhnikov. On uniqueness of probability solutions of the fokker-planck-kolmogorov equation. *Sbornik: Mathematics*, 212(6):745, 2021.
- [7] Ricky TQ Chen, Yulia Rubanova, Jesse Bettencourt, and David K Duvenaud. Neural ordinary differential equations. *Advances in neural information processing systems*, 31, 2018.
- [8] Valentin De Bortoli, James Thornton, Jeremy Heng, and Arnaud Doucet. Diffusion schrödinger bridge with applications to score-based generative modeling. *Advances in Neural Information Processing Systems*, 34:17695–17709, 2021.
- [9] Prafulla Dhariwal and Alexander Nichol. Diffusion models beat gans on image synthesis. *Advances in Neural Information Processing Systems*, 34:8780–8794, 2021.
- [10] Laurent Dinh, David Krueger, and Yoshua Bengio. Nice: Non-linear independent components estimation. *arXiv preprint arXiv:1410.8516*, 2014.
- [11] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. *Advances in neural information processing systems*, 27, 2014.
- [12] Will Grathwohl, Ricky TQ Chen, Jesse Bettencourt, Ilya Sutskever, and David Duvenaud. Ffjord: Free-form continuous dynamics for scalable reversible generative models. *arXiv preprint arXiv:1810.01367*, 2018.
- [13] Nikita Gushchin, Alexander Kolesov, Alexander Korotin, Dmitry Vetrov, and Evgeny Burnaev. Entropic neural optimal transport via diffusion processes. *arXiv preprint arXiv:2211.01156*, 2022.
- [14] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 33:6840–6851, 2020.
- [15] Jonathan Ho and Tim Salimans. Classifier-free diffusion guidance. *arXiv preprint arXiv:2207.12598*, 2022.
- [16] Michael F Hutchinson. A stochastic estimator of the trace of the influence matrix for laplacian smoothing splines. *Communications in Statistics-Simulation and Computation*, 18(3):1059–1076, 1989.
- [17] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [18] Alexander Korotin, Daniil Selikhanovych, and Evgeny Burnaev. Neural optimal transport. *arXiv preprint arXiv:2201.12220*, 2022.
- [19] Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.

- [20] Guan-Horng Liu, Arash Vahdat, De-An Huang, Evangelos A Theodorou, Weili Nie, and Anima Anandkumar.  $\mathbb{Q}$  sb: Image-to-image schrödinger bridge. *arXiv preprint arXiv:2302.05872*, 2023.
- [21] Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and transfer data with rectified flow. *arXiv preprint arXiv:2209.03003*, 2022.
- [22] Danilo Rezende and Shakir Mohamed. Variational inference with normalizing flows. In *International conference on machine learning*, pages 1530–1538. PMLR, 2015.
- [23] Yuyang Shi, Valentin De Bortoli, Andrew Campbell, and Arnaud Doucet. Diffusion schrödinger bridge matching. *arXiv preprint arXiv:2303.16852*, 2023.
- [24] Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. *Advances in neural information processing systems*, 32, 2019.
- [25] Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. *arXiv preprint arXiv:2011.13456*, 2020.
- [26] Alexander Tong, Nikolay Malkin, Guillaume Huguet, Yanlei Zhang, Jarrid Rector-Brooks, Kilian Fatras, Guy Wolf, and Yoshua Bengio. Conditional flow matching: Simulation-free dynamic optimal transport. *arXiv preprint arXiv:2302.00482*, 2023.
- [27] Max Welling and Yee W Teh. Bayesian learning via stochastic gradient langevin dynamics. In *Proceedings of the 28th international conference on machine learning (ICML-11)*, pages 681–688, 2011.